

ImputeZScore.jl

September 12, 2017

1 Imputation of association statistics in GWAS

1.1 Background

In a traditional GWAS, one collects genotype data at a small subset of SNPs over some individuals, then imputes genotypes across the entire genome, and finally computes association statistics (e.g. Z-scores) for each genotyped and imputed SNPs. This procedure can take tremendous amount of time as genotype imputation is computationally extensive.

Here, we impute Z-scores of ungenotyped SNPs directly from Z-scores of genotyped SNPs, without first performing genotype imputation, saving hundreds of hours of CPU time. The idea behind this approach is that Z-scores of genotyped and ungenotyped SNPs follow a multivariate normal distribution with LD matrix, which can be estimated from a reference panel, as the covariance structure -- one can impute the Z-scores of ungenotyped SNPs as the expectation of Z-scores of ungenotyped SNPs conditional on the Z-scores of genotyped SNPs.

In detail, let $Z = (Z_t, Z_u)$ be the Z-score vector partitioned into two components, genotyped (Z_t) and ungenotyped (Z_u). It has been previously shown that Z has the following distribution,

$$\begin{bmatrix} Z_t \\ Z_u \end{bmatrix} \sim MVN \left(\begin{bmatrix} \Lambda_t \\ \Lambda_u \end{bmatrix}, \begin{bmatrix} \Sigma_{tt} & \Sigma_{tu} \\ \Sigma_{ut} & \Sigma_{uu} \end{bmatrix} \right),$$

where $\Lambda = (\Lambda_t, \Lambda_u)$ is the non-centrality parameter, Σ_{tt} the LD between genotyped SNPs, Σ_{tu} the LD between genotyped and ungenotyped SNPs, Σ_{uu} the LD between ungenotyped SNPs.

The conditional expectation of Z_u given Z_t is then

$$Z_u | Z_t \sim MVN \left(\Lambda_u + \Sigma_{ut} \Sigma_{tt}^{-1} Z_t, \Sigma_{uu} - \Sigma_{ut} \Sigma_{tt}^{-1} \Sigma_{tu} \right).$$

We impute the Z-scores of ungenotyped SNPs as $\hat{Z}_u = E[Z_u | Z_t] = \Sigma_{ut} \Sigma_{tt}^{-1} Z_t$ under the null assumption that $\Lambda_u = 0$. Let $W = \Sigma_{ut} \Sigma_{tt}^{-1}$. This can be viewed as the weights on Z-scores of genotyped SNPs in the imputation of Z-scores of ungenotyped SNPs. Then

$$\hat{Z}_u \sim MVN(0, \Sigma_{ut} \Sigma_{tt}^{-1} \Sigma_{tu}),$$

where each entry $\hat{Z}_{u,i}$ of \hat{Z}_u follows

$$\hat{Z}_{u,i} \sim N(0, \Sigma_{ut,i*} \Sigma_{tt}^{-1} \Sigma_{tu,*i}).$$

Here, $\Sigma_{ut,i*}$ denotes the i -th row of Σ_{ut} and $\Sigma_{tu,*i}$ the i -th column of Σ_{tu} . To obtain a associations statistics that has mean 0 and variance 1, we standardize $\hat{Z}_{u,i}$ by $\sqrt{\Sigma_{ut,i*}\Sigma_{tt}^{-1}\Sigma_{tu,*i}}$. More specifically, the final imputed association statistics of each SNP is

$$\hat{Z}_{imp,i} = \frac{\hat{Z}_{u,i}}{\sqrt{\Sigma_{ut,i*}\Sigma_{tt}^{-1}\Sigma_{tu,*i}}} \sim N(0,1).$$

In practice, inverting a large matrix can be time-consuming. Instead, we adopt a window-based approach, i.e. we impute Z-scores of ungenotyped SNPs one window at a time.

2 Example

```
In [1]: # load in required packages
include("../src/ImputeZScore.jl")
using DataFrames, ImputeZScore, SnpArrays

In [2]: # read in z-scores of genotyped SNPs on chromosome 22 as a data frame
zsc_t = readtable("./hdl_chr22_typed.zsc", separator=' ')
```

```
Out[2]: 3000E5 DataFrames.DataFrame
  Row  rsID      pos  AO  A1  Z
1   "rs5746647"  17057138 "G" "T" 1.9927
2   "rs5747968"  17067504 "G" "T" 3.07937
3   "rs2236639"  17072483 "A" "G" 2.16346
4   "rs5746679"  17080378 "A" "G" 2.58537
5   "rs11089263" 17087656 "C" "A" 0.5
6   "rs2096537"  17094749 "A" "C" -0.291139
7   "rs4819849"  17152611 "A" "G" 2.12903
8   "rs2845379"  17202602 "T" "C" -0.0188679
9   "rs2845346"  17214252 "C" "T" -0.144928
10  "rs1807512"  17221495 "T" "C" 0.190476
11  "rs5748593"  17227461 "T" "C" -0.028169

2989 "rs6010023"  51028202 "C" "T" -0.731959
2990 "rs140510"   51052379 "T" "C" 0.0384615
2991 "rs131729"   51053719 "T" "C" 0.581818
2992 "rs3091400"  51059118 "G" "T" -0.344828
2993 "rs8142033"  51062832 "G" "A" 1.18462
2994 "rs9616812"  51105556 "C" "T" -1.625
2995 "rs2341009"  51133580 "C" "A" -1.13433
2996 "rs736334"   51139178 "C" "T" 0.278689
2997 "rs6010063"  51156933 "A" "G" 0.44
2998 "rs8137951"  51165664 "G" "A" -0.368421
2999 "rs3810648"  51175626 "A" "G" 0.415842
3000 "rs2238837"  51212875 "A" "C" 0.448718
```

```
In [3]: # there are only 3000 genotyped SNPs
        println(size(zsc_t))
```

```
(3000,5)
```

```
In [4]: # read in reference panel from 1000 genomes project for chromosome 22
        refpanel = SnpData("./1000G.EUR.22");
```

```
In [5]: # the reference panel has 17,489 SNPs on chromosome 22
        println(size(refpanel.snpmatrix))
```

```
(489,17489)
```

```
In [6]: # perform imputation using the method described above
        # this will take less than 20 seconds
        @time zsc_imp = impute_zscore(zsc_t, refpanel)
```

```
12.159418 seconds (56.25 M allocations: 1.758 GB, 3.29% gc time)
```

```
Out[6]: 10268x6 DataFrames.DataFrame
```

Row	rsID	pos	A0	A1	Z	r2pred
1	"rs2379981"	17030792	"G"	"A"	1.92311	0.705895
2	"rs4535153"	17031072	"C"	"T"	1.92311	0.705895
3	"rs9605903"	17054720	"C"	"T"	2.07014	0.740814
4	"rs5747988"	17073066	"A"	"G"	2.14202	0.888847
5	"rs5746664"	17074622	"A"	"C"	2.23801	0.914545
6	"rs2070501"	17084609	"A"	"G"	0.345276	0.915018
7	"rs16984366"	17096864	"C"	"T"	-1.94536	0.721418
8	"rs8137637"	17103717	"G"	"T"	-1.98102	0.729256
9	"rs4410381"	17107266	"A"	"G"	-2.10321	0.73318
10	"rs5993671"	17116398	"G"	"T"	-0.246478	0.894049
11	"rs5992472"	17132490	"G"	"A"	-0.269963	0.886824
10257	"rs762672"	51064818	"T"	"C"	-0.0281581	0.859197
10258	"rs10854884"	51101899	"A"	"C"	1.52138	0.818583
10259	"rs8138460"	51103692	"G"	"A"	1.68241	0.837635
10260	"rs9616906"	51104680	"A"	"G"	1.64511	0.851399
10261	"rs9628185"	51109992	"C"	"T"	1.55364	0.908213
10262	"rs9616915"	51117580	"T"	"C"	-1.57798	0.820481
10263	"rs9616816"	51123505	"A"	"G"	-0.350208	0.805102
10264	"rs739365"	51140316	"T"	"C"	-0.516951	0.679435
10265	"rs10451"	51162059	"A"	"G"	0.18436	0.904821
10266	"rs715586"	51163138	"T"	"C"	-0.778781	0.658237
10267	"rs2285395"	51178090	"A"	"G"	-0.374061	0.855836
10268	"rs3865766"	51186228	"T"	"C"	-0.473425	0.735647

```
In [7]: # the imputed data set has much more SNPs
# the last column r2pred is a measure of imputation accuracy
# note that this module filters out poorly imputed SNPs (r2pred < 0.6) by default
println(size(zsc_imp))
```

(10268,6)

```
In [8]: # now let's compare the imputed z-scores with the z-scores obtained by first
# performing a genotype imputation
zsc_full = readtable("./hdl_chr22_full.zsc", separator=' ');

# this removes SNPs with duplicated SNP ID and position and makes the
# sign consistent
snp_legend = DataFrame(rsID = refpanel.snpid, pos = refpanel.basepairs,
    A0 = refpanel.allele1, A1 = refpanel.allele2)
filter_input!(zsc_full, snp_legend)
```

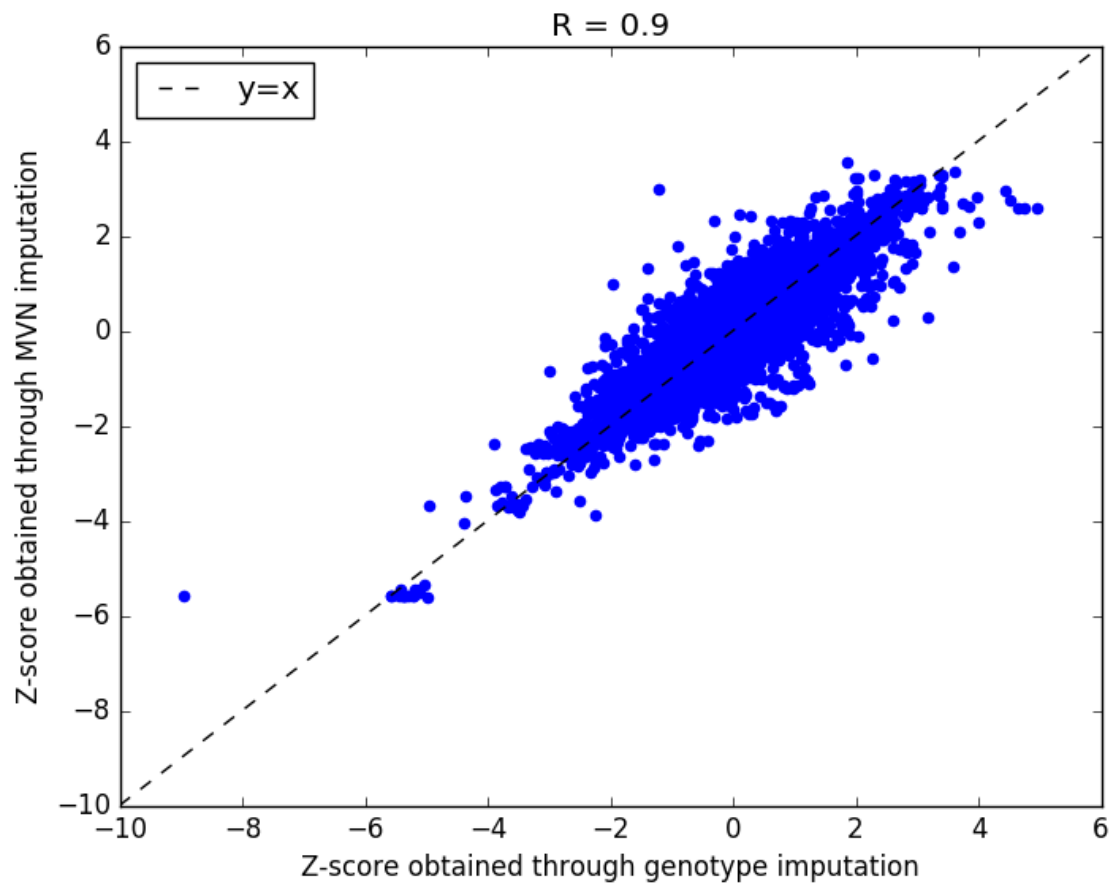
```
In [9]: # match z-scores based on SNP ID
zsc_matched = join(zsc_full, zsc_imp, on=:rsID)
println(zsc_matched)
```

7444E10 DataFrames.DataFrame

Row	rsID	pos	A0	A1	Z	pos_1	A0_1
1	"rs1000427"	36890105	"G"	"A"	-0.697368	36890105	"A"
2	"rs1000815"	26831077	"A"	"G"	0.166667	26831077	"A"
3	"rs1001021"	26403599	"G"	"A"	-0.290598	26403599	"A"
4	"rs1001213"	34131736	"G"	"A"	1.66346	34131736	"A"
5	"rs1001586"	42670293	"G"	"T"	0.516667	42670293	"T"
6	"rs1001587"	42670111	"C"	"T"	0.516667	42670111	"T"
7	"rs1001794"	32850930	"C"	"T"	-0.163636	32850930	"T"
8	"rs1001896"	19032215	"G"	"A"	-0.580645	19032215	"A"
9	"rs1002048"	34253393	"G"	"T"	1.2	34253393	"G"
10	"rs1002189"	30771458	"T"	"C"	-1.69643	30771458	"C"
11	"rs1003480"	31346752	"A"	"G"	-0.1875	31346752	"A"
7433	"rs9941935"	22015144	"A"	"G"	2.8	22015144	"G"
7434	"rs9941962"	40172198	"A"	"G"	-0.525641	40172198	"G"
7435	"rs9941971"	26051000	"T"	"C"	1.15238	26051000	"C"
7436	"rs9956"	32015450	"T"	"G"	1.31481	32015450	"G"
7437	"rs9967"	18211205	"T"	"C"	1.08333	18211205	"C"
7438	"rs997120"	33108536	"C"	"T"	-0.102564	33108536	"T"
7439	"rs997379"	35274298	"G"	"T"	0.270833	35274298	"G"
7440	"rs9983"	30423744	"G"	"A"	0.461538	30423744	"A"
7441	"rs998482"	41108135	"G"	"A"	-0.179104	41108135	"A"
7442	"rs9985182"	45539841	"G"	"T"	-0.0392157	45539841	"T"
7443	"rs999540"	37121101	"G"	"A"	0.174603	37121101	"A"
7444	"rs9997"	20796175	"T"	"C"	0.7	20796175	"C"

Row	A1_1	Z_1	r2pred
1	"G"	-0.728381	0.924804
2	"G"	-0.210362	0.942315
3	"G"	-0.476867	0.903856
4	"G"	1.24971	0.911577
5	"G"	0.440369	0.979745
6	"C"	0.473304	0.976182
7	"C"	0.086738	0.897738
8	"G"	-0.530442	0.894444
9	"T"	1.20897	0.905532
10	"T"	-1.8985	0.963061
11	"G"	-0.0941326	0.942817
7433	"A"	2.83774	0.817645
7434	"A"	-0.706095	0.974607
7435	"T"	1.11075	0.959307
7436	"T"	0.948379	0.956148
7437	"T"	1.30208	0.914667
7438	"C"	-1.09414	0.809205
7439	"T"	0.451928	0.957169
7440	"G"	0.416925	0.976849
7441	"G"	-0.0895031	0.971362
7442	"G"	-0.0544267	0.859937
7443	"G"	-0.00408845	0.79141
7444	"T"	0.850073	0.952408

```
In [11]: # make scatter plot comparing the z-scores obtained through different methods
         using PyPlot
         scatter(zsc_matched[:Z], zsc_matched[:Z_1], color="blue")
         plot([-10, 6], [-10, 6], color="black", linestyle="--", label="y=x")
         legend(loc=2)
         xlim([-10, 6]); ylim([-10, 6])
         xlabel("Z-score obtained through genotype imputation")
         ylabel("Z-score obtained through MVN imputation")
         title("R = " * string(round(cor(zsc_matched[:Z], zsc_matched[:Z_1]),2)))
```



Out[11]: PyObject <matplotlib.text.Text object at 0x7f90ed6fbbd0>

In []: